

# Modeling class probabilities via logistic regression

- Logistic regression intuition and conditional probabilities
- Learning the weights of the logistic cost function
- Converting an Adaline implementation into an algorithm for logistic regression
- Training a logistic regression model with scikit-learn
- Tackling overfitting via regularization

- Logistic regression과 conditional probabilities

- odds ratio  $\frac{p}{(1-p)}$



- logit function  $\text{logit}(p) = \log \frac{p}{(1-p)}$

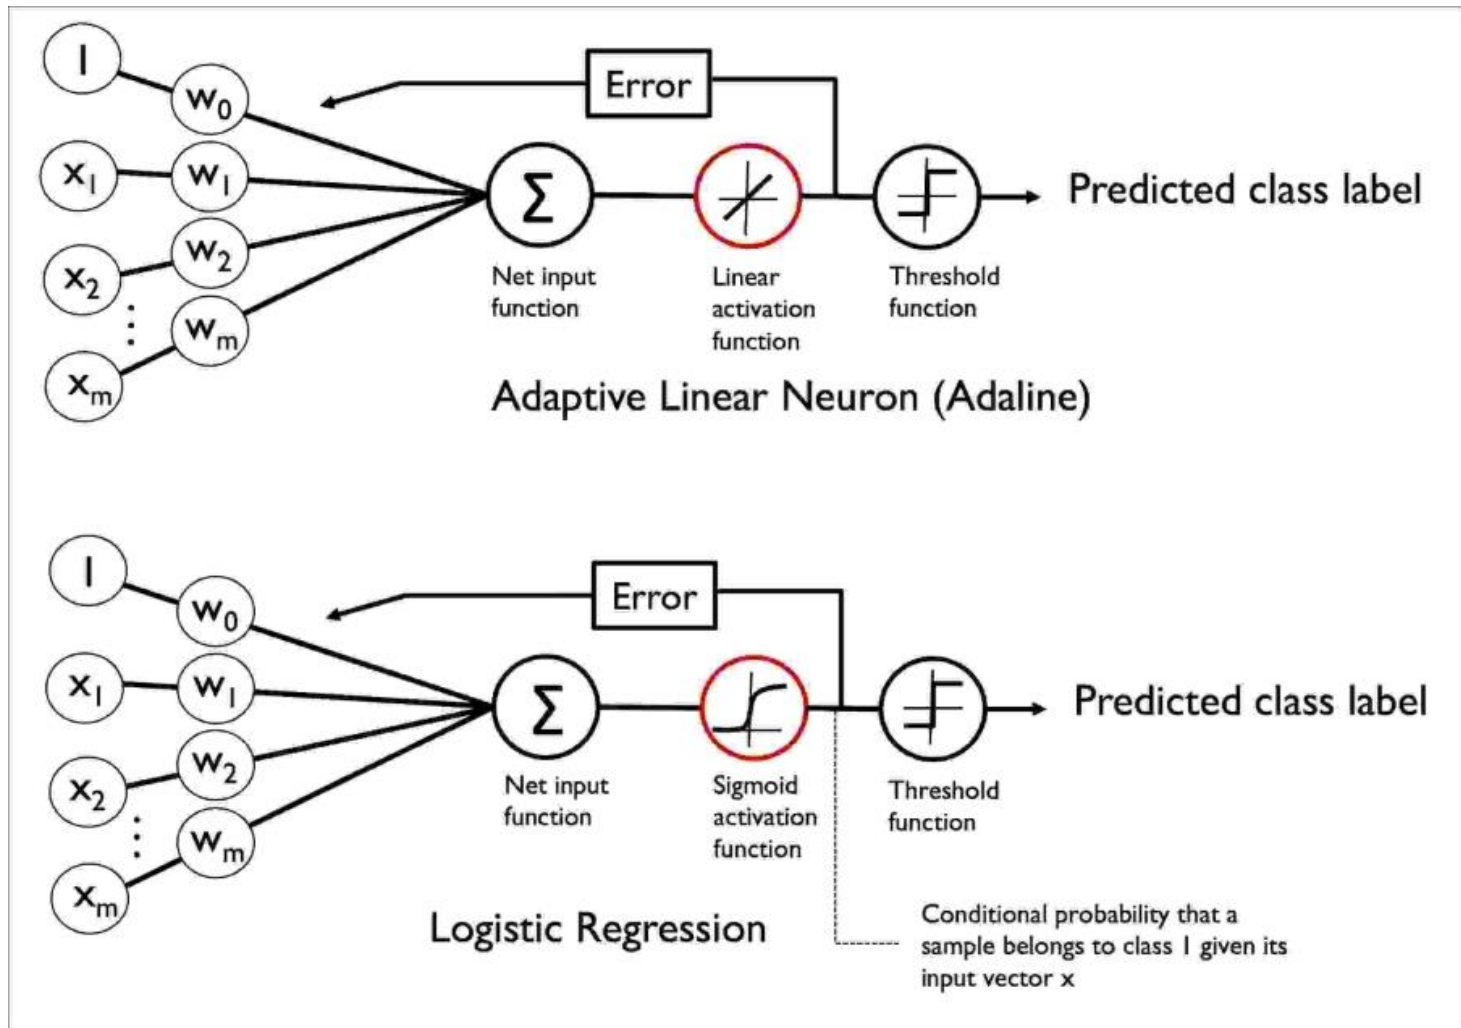
$$\text{logit}(p(y=1|\mathbf{x})) = w_0x_0 + w_1x_1 + \dots + w_mx_m = \sum_{i=0}^n w_ix_i = \mathbf{w}^T \mathbf{x}$$

- logistic function

- sigmoid function  $\phi(z) = \frac{1}{1+e^{-z}}$

- (단,  $z$  is a linear combination of weights & sample features)





$$\hat{y} = \begin{cases} 1 & \text{if } \phi(z) \geq 0.5 \\ 0 & \text{otherwise} \end{cases}$$

$$\hat{y} = \begin{cases} 1 & \text{if } z \geq 0.0 \\ 0 & \text{otherwise} \end{cases}$$

- Logistic cost function의 weights 학습

$$J(\mathbf{w}) = \sum_i \frac{1}{2} \left( \phi(z^{(i)}) - y^{(i)} \right)^2$$

- Define likelihood L to maximize

- 단, assume independence of individual samples

$$L(\mathbf{w}) = P(\mathbf{y} | \mathbf{x}; \mathbf{w}) = \prod_{i=1}^n P(y^{(i)} | x^{(i)}; \mathbf{w}) = \prod_{i=1}^n \left( \phi(z^{(i)}) \right)^{y^{(i)}} \left( 1 - \phi(z^{(i)}) \right)^{1-y^{(i)}}$$

- log-likelihood function


$$l(\mathbf{w}) = \log L(\mathbf{w}) = \sum_{i=1}^n \log \left( \phi(z^{(i)}) \right) + (1 - y^{(i)}) \log \left( 1 - \phi(z^{(i)}) \right)$$

# Training and Cost Function

- Cost function
  - $-\log(t)$  grows large when  $t$  approach 0  $\rightarrow$  cost will be large if ...
  - Vice versa

$$J(\phi(z), y; \mathbf{w}) = \begin{cases} -\log(\phi(z)) & \text{if } y = 1 \\ -\log(1 - \phi(z)) & \text{if } y = 0 \end{cases}$$

- cost function over whole training



$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(\hat{p}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{p}^{(i)})]$$

- Converting an Adaline implementation into an algorithm for logistic regression

$$J(\mathbf{w}) = -\sum_i y^{(i)} \log(\phi(z^{(i)})) + (1 - y^{(i)}) \log(1 - \phi(z^{(i)}))$$

```
class LogisticRegressionGD(object):
```

```
    """
```

```
    Parameters
```

```
    -----
```

```
    eta : float : Learning rate (between 0.0 and 1.0)
```

```
    n_iter : int : Passes over the training dataset.
```

```
    random_state : int : Random number generator seed for random weight initialization.
```

```
    Attributes
```

```
    -----
```

```
    w_ : 1d-array : Weights after fitting.
```

```
    cost_ : list : Sum-of-squares cost function value in each epoch.
```

```
    """
```

```
    def __init__(self, eta=0.05, n_iter=100, random_state=1):
```

```
        self.eta = eta
```

```
        self.n_iter = n_iter
```

```
        self.random_state = random_state
```

```
    def fit(self, X, y):
```

```
        """ Fit training data.
```

```
        Parameters
```

```
        -----
```

```
        X : {array-like}, shape = [n_samples, n_features]
```

```
        Training vectors, where n_samples is the number of samples and  
        n_features is the number of features.
```

```
        y : array-like, shape = [n_samples] : Target values.
```

```
    Returns
```

```
    -----
```

```
    self : object
```

```
    """
```

```

rgen = np.random.RandomState(self.random_state)
self.w_ = rgen.normal(loc=0.0, scale=0.01, size=1 + X.shape[1])
self.cost_ = []

for i in range(self.n_iter):
    net_input = self.net_input(X)
    output = self.activation(net_input)
    errors = (y - output)
    self.w_[1:] += self.eta * X.T.dot(errors)
    self.w_[0] += self.eta * errors.sum()

    # note that we compute the logistic `cost` now
    # instead of the sum of squared errors cost
    cost = (-y.dot(np.log(output)) - ((1 - y).dot(np.log(1 - output))))
    self.cost_.append(cost)
return self

def net_input(self, X):
    """Calculate net input"""
    return np.dot(X, self.w_[1:]) + self.w_[0]

def activation(self, z):
    """Compute logistic sigmoid activation"""
    return 1. / (1. + np.exp(-np.clip(z, -250, 250)))

def predict(self, X):
    """Return class label after unit step"""
    return np.where(self.net_input(X) >= 0.0, 1, 0)
    # equivalent to:
    # return np.where(self.activation(self.net_input(X)) >= 0.5, 1, 0)


```



- Training a logistic regression model with scikit-learn

```
>>> from sklearn.linear_model import LogisticRegression
>>> lr = LogisticRegression(C=1000.0, random_state=0)
>>> lr.fit(X_train_std, y_train)
>>> plot_decision_regions(X_combined_std,
...                       y_combined, classifier=lr,
...                       test_idx=range(105,150))
>>> plt.xlabel('petal length [standardized]')
>>> plt.ylabel('petal width [standardized]')
>>> plt.legend(loc='upper left')
>>> plt.show()
```

- Tackling overfitting via regularization
  - Overfitting vs. Underfitting



- [note]
  - Variance measures the consistency (or variability) of the model prediction for a particular sample instance if we would retrain the model multiple times
  - bias measures how far off the predictions are from the correct values in general;
    - measure of the systematic error that is not due to randomness.
    - 예: L2 regularization  $C = \frac{1}{\lambda}$

$$\frac{\lambda}{2} \|\mathbf{w}\|^2 = \frac{\lambda}{2} \sum_{j=1}^m w_j^2$$

$$J(\mathbf{w}) = C \left[ \sum_{i=1}^n \left( -\log(\phi(z^{(i)})) + (1 - y^{(i)}) (-\log(1 - \phi(z))) \right) \right] + \frac{1}{2} \|\mathbf{w}\|^2$$